



رواد مصر الرقمية

Amazon Customer Product Reviews



Table Of Contents

01	Introduction	06	Modeling Approaches
02	Team Members	07	Experiment Tracking
03	Dataset	08	Model Deployment
04	Exploratory Data Analysis	09	Results and Evaluation
05	Preprocessing and Feature Engineering	10	Conclusion

Introduction

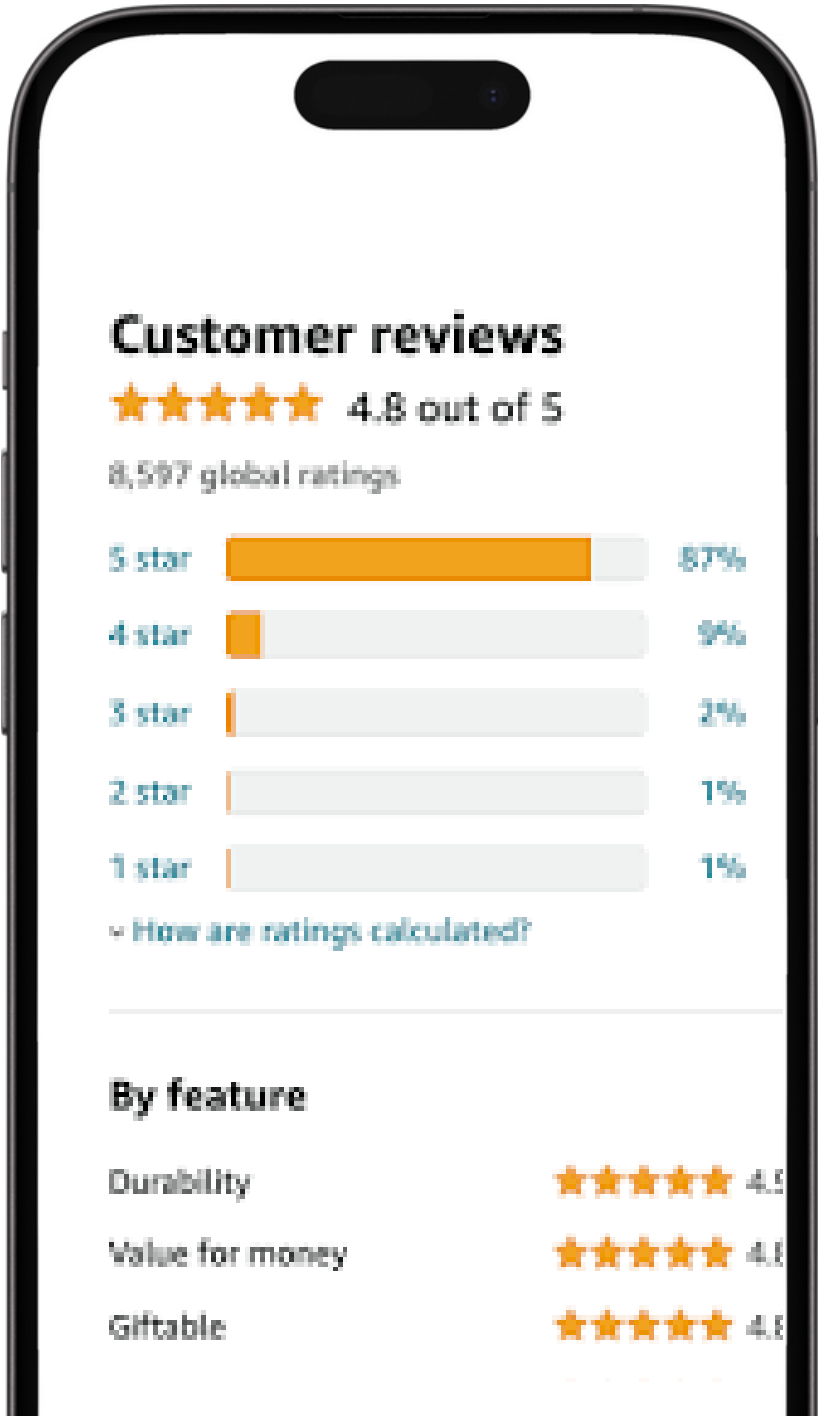
We performed sentiment analysis on Amazon customer reviews to classify them as positive or negative. The project involves using machine learning and deep learning models, including LSTM, to achieve high accuracy in predicting customer sentiment. Additionally, we implemented platforms like Hugging Face, MLflow, and Streamlit for model optimization, tracking, and deployment.



Team Members

- Mariam Mohamed
- Haidy Talaat
- Ahmed Emad
- Mostafa Omran
- Abdelrahman Tarek
- Abdelrahman Elsheikh
- Fatma Salah

DATASET



Dataset Overview


This dataset contains over 568,000 consumer reviews ratings scored 1-5 stars for a variety of Amazon products.

Data dictionary:



- **Id:** Unique identifier for each review
- **ProductId:** Amazon product ID
- **UserId:** Amazon user ID
- **ProfileName:** User's profile name
- **HelpfulnessNumerator:** Number of users who found the review helpful
- **HelpfulnessDenominator:** Total number of votes for the review
- **Score:** Review rating (1-5 stars)
- **Time:** Timestamp of the review
- **Summary:** Short summary of the review
- **Text:** Full review text

EXPLORATORY DATA ANALYSIS (EDA)

- 1. Score distribution**
- 2. Word Cloud**
- 3. Data Duplication**




Score Distribution



Key Observations:

- **Imbalanced data**
- **Fewer lower scores:** Ratings of 1, 2, and 3 are relatively infrequent, indicating a lower prevalence of negative or neutral sentiment.

Word Cloud



Key observation:

- **The word cloud visually represents the most frequently used terms in the dataset, highlighting words like "coffee," "tea," "taste," "good," "product," and "use," suggesting that these topics are central to the customer reviews.**

Data Duplication

```
# To check for duplicates  
df.duplicated().sum()
```


```
174779
```

Key observation:

- There are 174779 duplicated reviews.

PREPROCESSING AND FEATURE ENGINEERING

- 1. Dropping Duplicated Values**
- 2. Sampling Data**
- 3. Text Preprocessing**
- 4. Score Encoding**



Dropping Duplicated values

```
df.drop_duplicates(inplace=True)
```

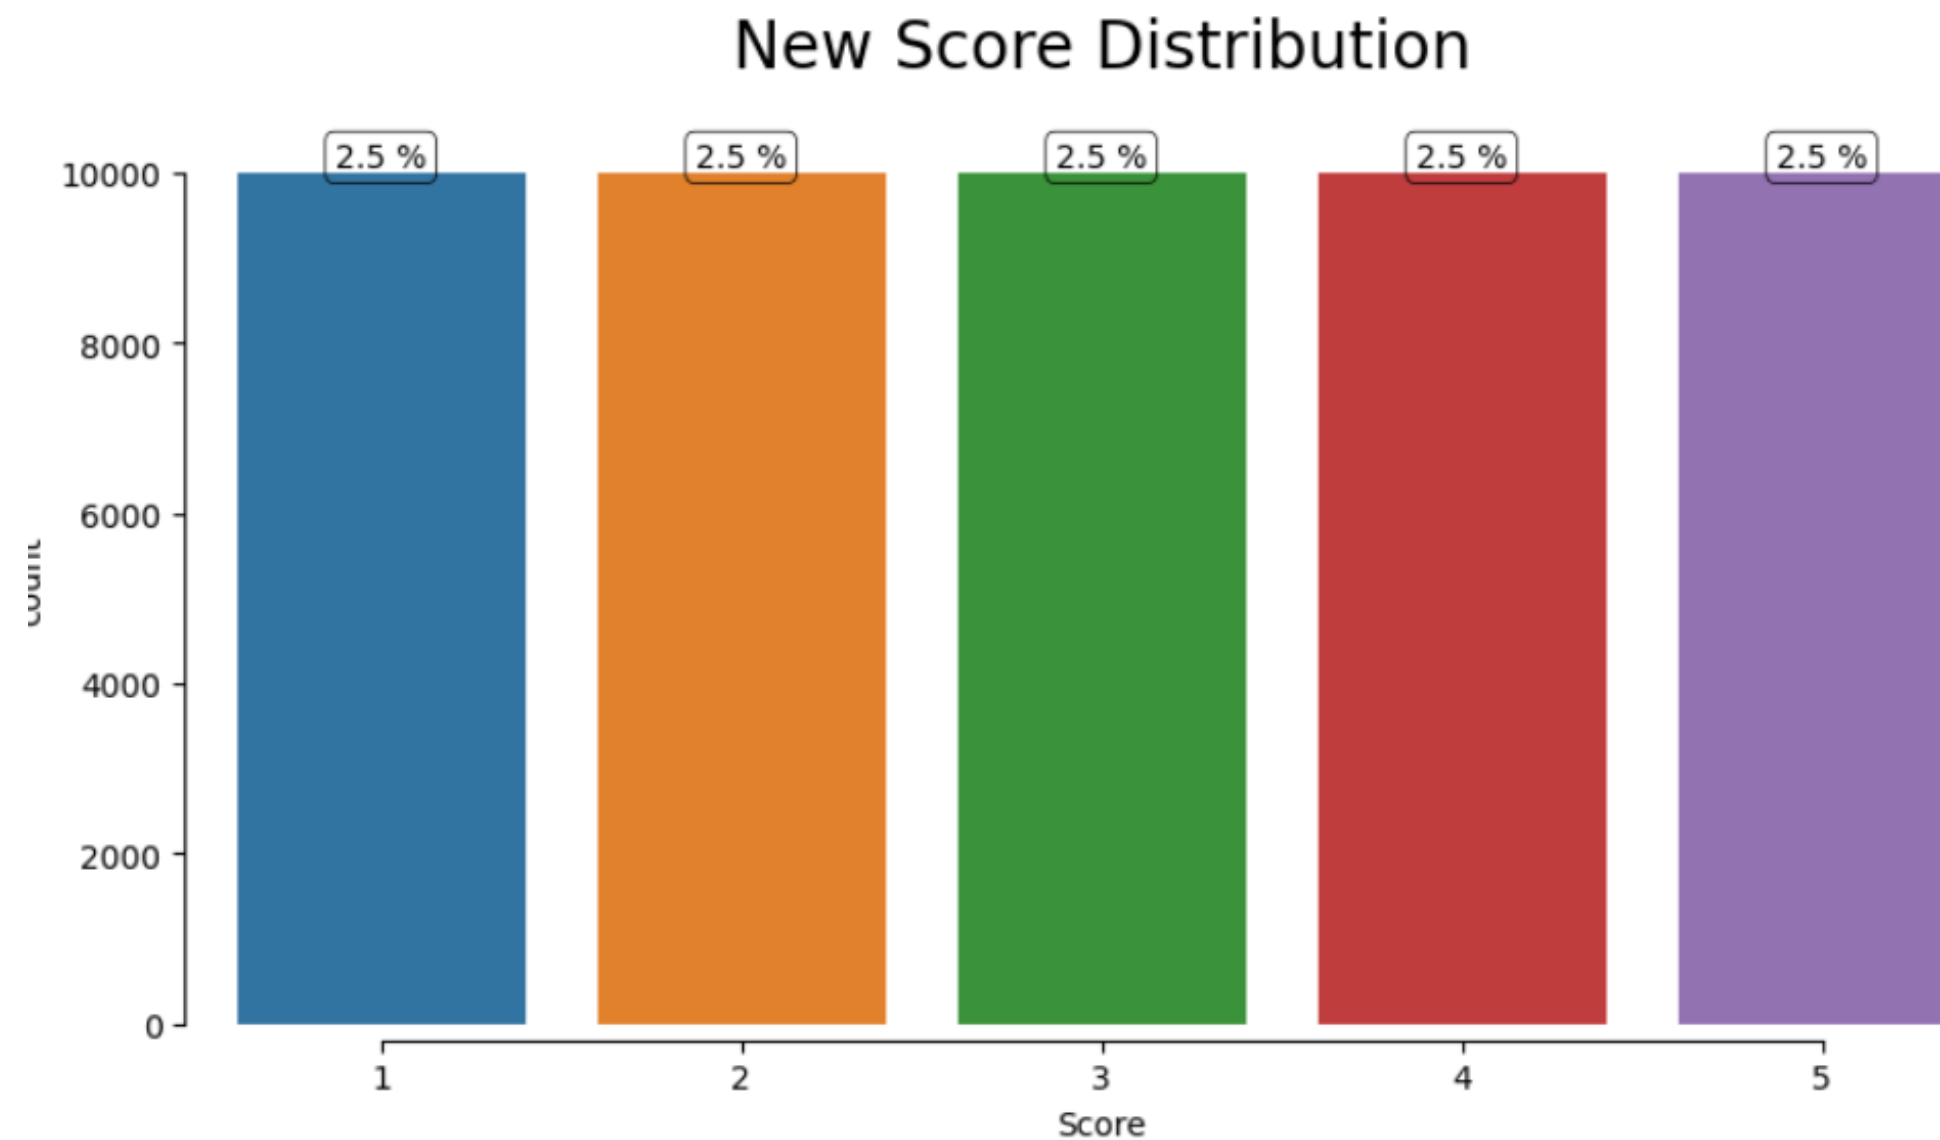
```
df.shape
```

```
(393675, 2)
```

To ensure the integrity of our dataset, we removed duplicate entries. The original dataset had 568454 and after dropping 174779 duplicates, the new dataset dimension is 393675. This step helps eliminate redundant data, improving the quality of the analysis.

As we don't need the other attributes, we drop all except the text and score.

Sampling Data



We implemented a balanced sampling strategy to address class imbalance by randomly selecting 10,000 reviews from each score category. This approach ensures equal representation of all score categories, facilitating more accurate sentiment analysis.

Text Preprocessing

```
def clean_text(text):  
    # 1. Convert to lower  
    txt=text.lower()  
  
    # 1. split to words  
    tokens=word_tokenize(text)  
  
    # 3. remove punctuation  
    tokens=[word for word in tokens if word not in string.punctuation]  
  
    # 4. Remove stopwords  
    tokens=[word for word in tokens if word not in stop_words]  
  
    # 5. Remove numbers  
    tokens=[word for word in tokens if not word.isdigit()]  
  
    # 6. Apply Stemming  
    tokens=[stemming.stem(word) for word in tokens]  
  
    # To return these single words back into one string  
    return ' '.join(tokens)
```


Text	cleaned_text
The wheat free brownie mix is not to my liking...	the wheat free browni mix like I wo n't reorde...
saltyness may be the "norm" for products like ...	salty may `` norm " product like doesnt mean ...
I purchased this product from Otto's because t...	I purchas product otto 's offer via amazon pri...
Switch to this food and my dog became very sic...	switch food dog becam sick Go onlin lookup dog...
please do not take this note as an attack on y...	pleas take note attack product It snoodl hate ...
...	...
I bought this for Halloween and I had SOOOOOO ...	I bought halloween I soooooo much candi left I...
My two cockers love this formula of chicken an...	My two cocker love formula chicken sweet potat...
Ive hunted high and low and have tried every t...	ive hunt high low tri everi type licoric toffe...
We have a rapidly growing 7 month Labradoodle ...	We rapidli grow month labradoodl puppi initi b...
we received this coffee yesterday, and have to...	receiv coffe yesterday tell love espresso n't ...

Using NLTK, the text preprocessing function performs several steps: converting to lowercase, tokenizing words, removing punctuation and stop words, excluding numbers, and applying stemming. This prepares the text for analysis, simplifying it for consistent feature extraction in machine learning models.

Score Encoding

```
new_df['Score'] = new_df['Score'].apply(lambda x: 1 if x >=3 else 0)  
# 1 --> Good  
#0 --> Bad
```

To improve model performance, we simplified the score classification. Initially, the 5-class classification (very bad, bad, good, very good, excellent) achieved an accuracy of only 47.9% using Logistic Regression. To address this, we applied binary classification: scores of 4 and 5 were labeled as Positive (1), and scores of 1, 2, and 3 as Negative (0). This binary approach better suited the sentiment analysis task.



MODELING APPROACHES

- 1. Traditional models (Logistic Regression, Naive Bayes, SVM)**
- 2. Deep Learning (LSTM)**
- 3. Hugging Face**




Logistic Regression

```
pipe = Pipeline(  
    [  
        ('vec', CountVectorizer(stop_words= "english")),  
        ('tfidf', TfidfTransformer()),  
        ('classifier', LogisticRegression()),  
    ]  
)
```


Training accuracy: 0.8317

Test accuracy: 0.7953



Training accuracy: 0.8591315789473685

Test accuracy: 0.8229473684210526



Logistic Regression was implemented in two configurations: the left side shows results using only review text, while the right side includes both text and summary. The pipeline outlines the key steps in model training. Combining text with the summary led to a noticeable improvement in performance, boosting accuracy by 3%, from 79.53% to 82.29%.

Naive Bayes

```
naive_bayes_pipeline = Pipeline([
    ('vec', CountVectorizer(stop_words='english')),
    ('tfidf', TfidfTransformer()),
    ('classifier', MultinomialNB())
])
```


	precision	recall	f1-score	support		precision	recall	f1-score	support
					0	0.7987	0.7369	0.7666	3748
0	0.8427	0.3717	0.5158	4036	1	0.8368	0.8790	0.8574	5752
1	0.6915	0.9531	0.8015	5964					
accuracy			0.7184	10000	accuracy			0.8229	9500
macro avg	0.7671	0.6624	0.6586	10000	macro avg	0.8178	0.8080	0.8120	9500
weighted avg	0.7525	0.7184	0.6862	10000	weighted avg	0.8218	0.8229	0.8216	9500

As shown, there is a significant difference between the two results. When using only the review text, Naive Bayes achieved 71.84% accuracy. However, by combining the text with the summary, the accuracy improved dramatically to 82.29%. This highlights the benefit of incorporating additional information for better model performance

SVC

```
svm_pipeline = Pipeline([
    ('vec', CountVectorizer(stop_words='english')),
    ('tfidf', TfidfTransformer()),
    ('classifier', SVC())
])
```

	precision	recall	f1-score	support
0	0.7757	0.6967	0.7341	4036
1	0.8080	0.8637	0.8349	5964
accuracy			0.7963	10000
macro avg	0.7919	0.7802	0.7845	10000
weighted avg	0.7950	0.7963	0.7942	10000



The SVM model was applied using only the review text, achieving an overall accuracy of 79.63%. The F1-score for the positive class (1) was 0.8349, while the negative class (0) had a lower F1-score of 0.7341. These results show that SVM performed well in identifying positive sentiments but was less effective for negative ones.


LSTM Model Architecture

The LSTM model consists of an embedding layer with a vocabulary size of 20,000 and output embedding dimension of 200, followed by an LSTM layer with 256 units, recurrent dropout (0.3), and standard dropout (0.3). A final Dense layer with a sigmoid activation function is used for binary classification, with L2 regularization to prevent overfitting. The model is designed to process padded sequences of review text for sentiment analysis


```
model=Sequential()  
model.add(Embedding(input_dim=voca_size,output_dim=embedding_size,input_length=max_len))  
model.add(LSTM(256,recurrent_dropout=0.3,dropout=0.3))  
model.add(Dropout(0.5))  
model.add(Dense(1,activation='sigmoid',kernel_regularizer=regularizers.l2(0.01)))
```

```
# Compile the model  
import tensorflow as tf  
#op=tf.keras.optimizers.Adam(learning_rate=0.001)  
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])  
  
# Early stopping to prevent overfitting  
early_stopping = EarlyStopping(monitor='val_loss', patience=5)  
  
model_checkpoint = ModelCheckpoint(  
    'best_model1.keras', # File path where the model will be saved  
    monitor='val_loss', # Metric to monitor  
    save_best_only=True, # Save only the model with the best validation loss  
    mode='min', # 'min' because lower loss is better  
    verbose=1 # Verbosity mode  
)  
  
# Train the LSTM model  
history1=model.fit(X_train, y_train,  
    epochs=12,  
    batch_size=128,  
    validation_data=(X_test, y_test)  
    , callbacks=[early_stopping,model_checkpoint])
```

LSTM Accuracy




Final Test Accuracy: 78.75%



Final Test Accuracy: 77.48%

The LSTM model achieved a final test accuracy of 78.75% when trained using only the review text, showing a slight improvement in performance compared to the Text + Summary configuration.

MODEL DEPLOYMENT



1. Streamlit



Streamlit Deployment Overview

Choose your model:

- ☒ SVM
- ☐ Naive Bayes
- ☐ Logistic Regression
- ☐ LSTM

Customer Product Reviews Sentiment Analysis App

Enter text for sentiment analysis:

Type here...

Analyze Sentiment

Users can input a product review and choose from four models for sentiment classification: SVM, Naive Bayes, Logistic Regression and LSTM. then predicts whether the review is Positive or Negative, displaying the raw prediction output and a visual sentiment label.

Postive **Reviwes**

Choose your model:

- ☒ SVM
☐ Naive Bayes
☐ Logistic Regression
☐ LSTM

Customer Product Reviews Sentiment Analysis App

Enter text for sentiment analysis:

the product arrived very quickly, I recommend this website

Analyze Sentiment

Raw prediction output: [1]

Prediction: Positive ❤️

Choose your model:

- ☐ SVM
☐ Naive Bayes
☒ Logistic Regression
☐ LSTM

Customer Product Reviews Sentiment Analysis App

Enter text for sentiment analysis:

the product arrived very quickly, I recommend this website

Analyze Sentiment

Raw prediction output: [1]

Prediction: Positive ❤️

Choose your model:

- ☐ SVM
☒ Naive Bayes
☐ Logistic Regression
☐ LSTM

Customer Product Reviews Sentiment Analysis App

Enter text for sentiment analysis:

the product arrived very quickly, I recommend this website

Analyze Sentiment

Raw prediction output: [1]

Prediction: Positive ❤️

Choose your model:

- ☐ SVM
☐ Naive Bayes
☐ Logistic Regression
☒ LSTM

Customer Product Reviews Sentiment Analysis App

Enter text for sentiment analysis:

the product arrived very quickly, I recommend this website

Analyze Sentiment

Raw prediction output: [[1]]

Prediction: Positive ❤️

Negative Reveiws

Choose your model:

☒ SVM

☐ Naive Bayes

☐ Logistic Regression

☐ LSTM

Customer Product Reviews Sentiment Analysis App

Enter text for sentiment analysis:

It was terrible experience

Analyze Sentiment

Raw prediction output: [0]

Prediction: Negative 😞

Choose your model:

☐ SVM

☒ Naive Bayes

☐ Logistic Regression

☐ LSTM

Sentiment Analysis

Enter text for sentiment analysis:

it was bad experience

Analyze Sentiment

Raw prediction output: [0]

Prediction: Negative 😞

Choose your model:

☐ SVM

☐ Naive Bayes

☒ Logistic Regression

☐ LSTM

Sentiment Analysis


Enter text for sentiment analysis:

it was bad experience

Analyze Sentiment

Raw prediction output: [0]

Prediction: Negative 😞



EXPERIMENT TRACKING

1.MLflow


2.Azure



Our **MLflow**

Machine learning tracking with **MLflow**

Exploring innovative strategies
for sustainable growth and
competitive advantage.



ONLINE SHOPPING

Runs

Evaluation

Experimental

Traces

Experimental

☰

📈

🔍 metrics.rmse < 1 and params.model = "tree" ⓘ

Time created ▾

State: Active ▾

Datasets ▾

⌵ Sort: Created ▾

📄 Columns ▾

⋮

🔄

+ New run

📅 Group by ▾

						Metrics				Parameters		Tags
<input type="checkbox"/>	Run Name	Created ⌵	Duration	Source	Models	accuracy	auc	f1-score	train_accuracy	max_depth	n_estimators	clf
<input type="checkbox"/>	🔵 powerful-fowl-112	✅ 16 minutes ago	2.2min	🏠 mlflow_...	🔗 sklearn	0.75375	0.830450135...	0.75375	0.994375	10	250	xgboost
<input type="checkbox"/>	🔵 sneaky-ray-661	✅ 26 minutes ago	4.2min	🏠 mlflow_...	🔗 sklearn	0.7425	0.822427339...	0.7425	0.999921875	20	300	xgboost
<input type="checkbox"/>	🔴 whimsical-fish-275	✅ 40 minutes ago	3.7min	🏠 mlflow_...	🔗 sklearn	0.7390625	0.822898855...	0.7390625	0.999921875	12	400	xgboost
<input type="checkbox"/>	🟡 welcoming-zebra-486	✅ 3 hours ago	3.3min	🏠 mlflow_...	🔗 sklearn	0.7484375	0.831002515...	0.7484375	0.999921875	15	300	xgboost
<input type="checkbox"/>	🟢 colorful-shrimp-339	✅ 4 hours ago	3.8min	🏠 mlflow_...	🔗 sklearn	0.755	0.842836499...	0.755	1	15	300	xgboost
<input type="checkbox"/>	🔵 zealous-sloth-506	✅ 4 hours ago	4.9min	🏠 mlflow_...	🔗 sklearn	0.75125	0.834254683...	0.75125	1	20	350	xgboost
<input type="checkbox"/>	🟡 powerful-deer-391	✅ 4 hours ago	5.1min	🏠 mlflow_...	🔗 sklearn	0.756875	0.832760248...	0.756875	0.999921875	20	350	xgboost
<input type="checkbox"/>	🔵 incongruous-jay-718	✅ 5 hours ago	4.9min	🏠 mlflow_...	🔗 sklearn	0.7615	0.843707843...	0.7615	0.999875	15	350	xgboost

XGBOOST EXPERIMENT

detection_with_xgb >

Comparing 8 Runs from 1 Experiment

Visualizations

Parallel Coordinates Plot Scatter Plot Box Plot Contour Plot


Parameters:

max_depth X n_estimators X

Metrics:

accuracy X

Clear All



COMPARING XGBOOST

detection_with_xgb >

incongruous-jay-718

Register model

Overview Model metrics System metrics Artifacts

- ▼ XGBClassifier
 - ▼ XG_first_try
 - MLmodel
 - conda.yaml
 - model.pkl
 - python_env.yaml
 - requirements.txt
 - XG_first_try_conf_matrix.png
 - XG_first_try_roc_curve_micro.png

XG_first_try_conf_matrix.png 13.71KB

Path: file:///C:/Users/DELL/Downloads/New%20folder%20%282%29/mlruns/571923209411826566/8df460718ad44c88a48fcef16c6cb71d/artifacts/XG_first_try_conf_matri...



XGBOOST CONFUSION MATRIX

Runs Evaluation Experimental Traces Experimental

☰

📈

🔍 metrics.rmse < 1 and params.model = "tree" ⓘ

Time created ▾

State: Active ▾

Datasets ▾

⌵ Sort: Created ▾

📄 Columns ▾

⋮ ↺

+ New run

📄 Group by ▾

						Metrics			Parameters			
<input type="checkbox"/>	Run Name	Created ⌵	Duration	Source	Models	accuracy	auc	f1-score	train_accuracy	C	max_iter	solver
<input type="checkbox"/>	magnificent-mole-59	✅ 17 minutes ago	8.7s	🏠 mlflow_...	🔗 sklearn	0.765	0.850456143...	0.765	0.84140625	0.8	100	lbfgs
<input type="checkbox"/>	clean-gnu-318	✅ 24 minutes ago	12.6s	🏠 mlflow_...	🔗 sklearn	0.76	0.834588689...	0.76	0.863828125	1.5	100	lbfgs
<input type="checkbox"/>	honorable-ray-289	✅ 38 minutes ago	6.1s	🏠 mlflow_...	🔗 sklearn	0.7621875	0.846051555...	0.7621875	0.819453125	0.3	100	lbfgs
<input type="checkbox"/>	learned-lamb-491	✅ 3 hours ago	12.1s	🏠 mlflow_...	🔗 sklearn	0.7721875	0.858601602...	0.7721875	0.8684375	2.0	100	lbfgs
<input type="checkbox"/>	delicate-tern-56	❌ 4 hours ago	11.2s	🏠 mlflow_...	🔗 sklearn	0.7834375	0.859545805...	0.7834375	0.867421875	2.0	100	lbfgs
<input type="checkbox"/>	redolent-pig-590	❌ 4 hours ago	9.1s	🏠 mlflow_...	🔗 sklearn	0.7790625	0.857091347...	0.7790625	0.83703125	0.7	100	lbfgs
<input type="checkbox"/>	likeable-colt-275	❌ 4 hours ago	10.0s	🏠 mlflow_...	🔗 sklearn	0.7696875	0.844425667...	0.7696875	0.844453125	0.7	100	lbfgs
<input type="checkbox"/>	beautiful-trout-887	❌ 5 hours ago	37.2min	🏠 mlflow_...	🔗 sklearn	0.779	0.859319483...	0.779	0.847375	1.0	100	lbfgs

LOGISTIC REGRESSION

[detection_with_logistic_regression](#) >

Comparing 8 Runs from 1 Experiment

Visualizations

Parallel Coordinates Plot

Scatter Plot

Box Plot

Contour Plot


Parameters:

max_iter X C X

Metrics:

accuracy X

Clear All



COMPARING LOGISTIC REGRESSION

detection_with_logistic_regression >

delicate-tern-56



Register model

Overview

Model metrics

System metrics

Artifacts

LogisticRegression

lr_first_try

MLmodel

conda.yaml

model.pkl

python_env.yaml


requirements.txt

lr_first_try_conf_matrix.png

lr_first_try_conf_matrix.png 11.86KB



Path: file:///C:/Users/DELL/Downloads/New%20folder%20%282%29/mlruns/932000396248193146/927a412e425b40889f26e51990cb5e2e/artifacts/lr_first_try_conf_matri...




LOGISTIC REGRESSION

- ▼ SVC
 - ▼ svc_first_try
 - MLmodel
 - conda.yaml
 - model.pkl
 - python_env.yaml
 - requirements.txt
 - svc_first_try_conf_matrix.png**

svc_first_try_conf_matrix.png 12.01KB


Path: file:///C:/Users/DELL/Downloads/New%20folder%20%282%29/mlruns/384230930060874105/e1cd66212f424021adc5397d09cc231f/artifacts/svc_first_try_conf_matr... [Download](#) [Copy](#)



SVC CONFUSION MATRIX

🔍 Search metric charts

🔄 Refresh




SVC MATRIXS

Our **MLflow**

Deep learning tracking with **MLflow**

Exploring innovative strategies
for sustainable growth and
competitive advantage.



ONLINE SHOPPING

Reviews_detection

[Provide Feedback](#)[Add Description](#)[Share](#)[Runs](#)[Evaluation](#)[Experimental](#)[Traces](#)[Experimental](#)

Time created

State: Active

Datasets

[+ New run](#)

Sort: Created

Columns

Group by

				Metrics			Parameters		
<input type="checkbox"/>	Run Name	Created	Models	accuracy	auc	f1-score	LSTM_units	batch_size	dropout
<input type="checkbox"/>	bittersweet-fowl-456	49 minutes ago	tensorflow	0.76975	0.847651539...	0.76975	256	64	0.3
<input type="checkbox"/>	luminous-stork-477	1 hour ago	tensorflow	0.771625	0.850329635...	0.771625	256	128	0.3
<input type="checkbox"/>	redolent-shrew-927	1 hour ago	tensorflow	0.773125	0.859444590...	0.773125	256	256	0.1
<input type="checkbox"/>	resilient-tern-441	1 hour ago	tensorflow	0.81725	0.892760311...	0.81725	256	256	0.1
<input type="checkbox"/>	welcoming-kit-103	1 hour ago	tensorflow	0.813875	0.889818087...	0.813875	256	128	0.1
<input type="checkbox"/>	amazing-stag-298	1 hour ago	tensorflow	0.7945	0.870294161...	0.7945	256	128	0.3
<input type="checkbox"/>	upset-fly-446	1 hour ago	tensorflow	0.794625	0.873849110...	0.794625	256	128	0.1
<input type="checkbox"/>	fearless-skunk-890	2 hours ago	tensorflow	0.816125	0.888580353...	0.816125	256	128	0.1
<input type="checkbox"/>	zealous-croc-823	2 hours ago	tensorflow	0.80025	0.876022231...	0.80025	256	128	0.3
<input type="checkbox"/>	stylish-sheep-525	2 hours ago	tensorflow	0.799	0.878618575...	0.799	256	128	0.3

24 matching runs

LSTM EXPERIMENT

Comparing 14 Runs from 1 Experiment

Visualizations

Parallel Coordinates Plot Scatter Plot Box Plot Contour Plot

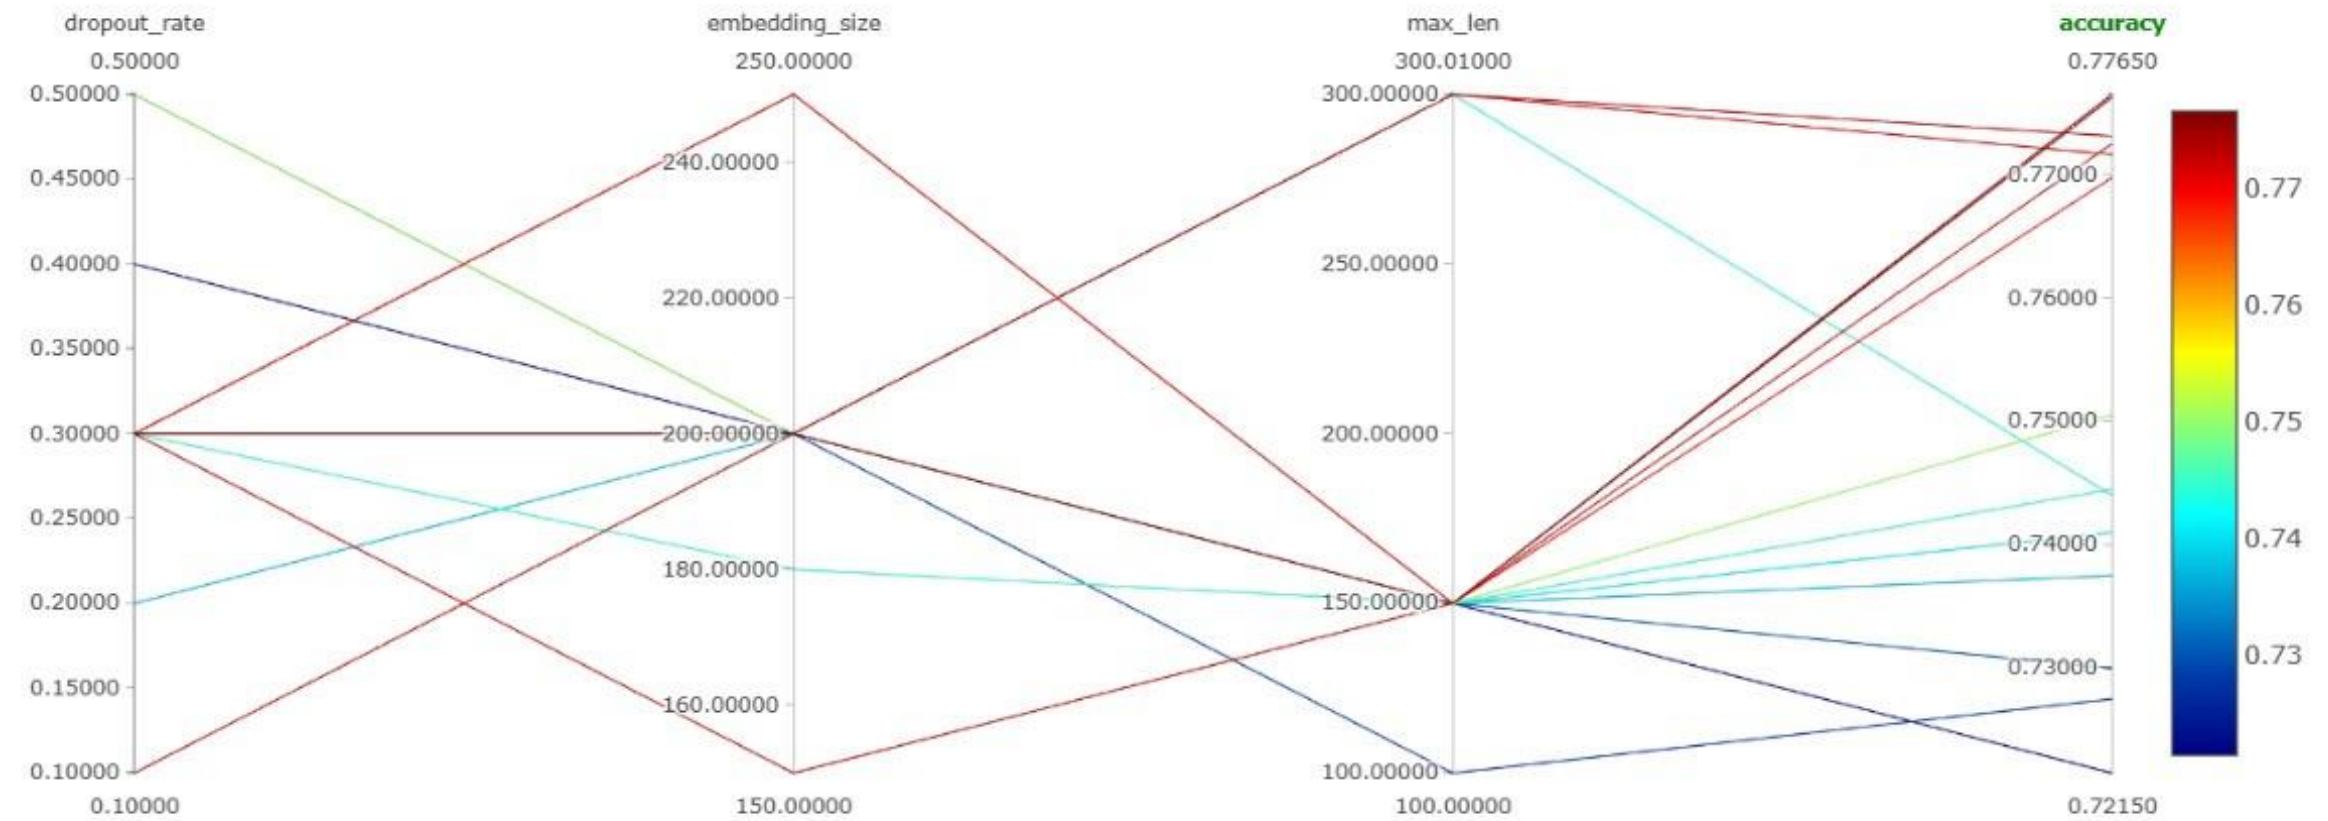
Parameters:

dropout_rate X embedding_size X max_len X

Metrics:

accuracy X


Clear All



COMPARING LSTM WITH TEXT

Overview Model metrics System metrics **Artifacts**

- Sequential
 - LSTM_first_try
 - data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - LSTM_first_try_conf_matrix.png**
 - LSTM_first_try_pr_curve.png
 - LSTM_first_try_roc_curve.png




LSTM

- ▼ Sequential
 - ▼ LSTM_first_try
 - ▶ data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - LSTM_first_try_conf_matrix.png
 - LSTM_first_try_pr_curve.png
 - LSTM_first_try_roc_curve.png

LSTM_first_try_roc_curve.png 26.55KB



Path: file:///c:/Users/DELL/Downloads/New%20folder%20%28%29/mlruns/940220449463371638/e5c59879170944f5bfeab26075d678f6/artifacts/LSTM_first_try_roc_curv...



LSTM ROC CURVE

Reviews_detection >

Comparing 10 Runs from 1 Experiment

Visualizations

Parallel Coordinates Plot

Scatter Plot

Box Plot

Contour Plot

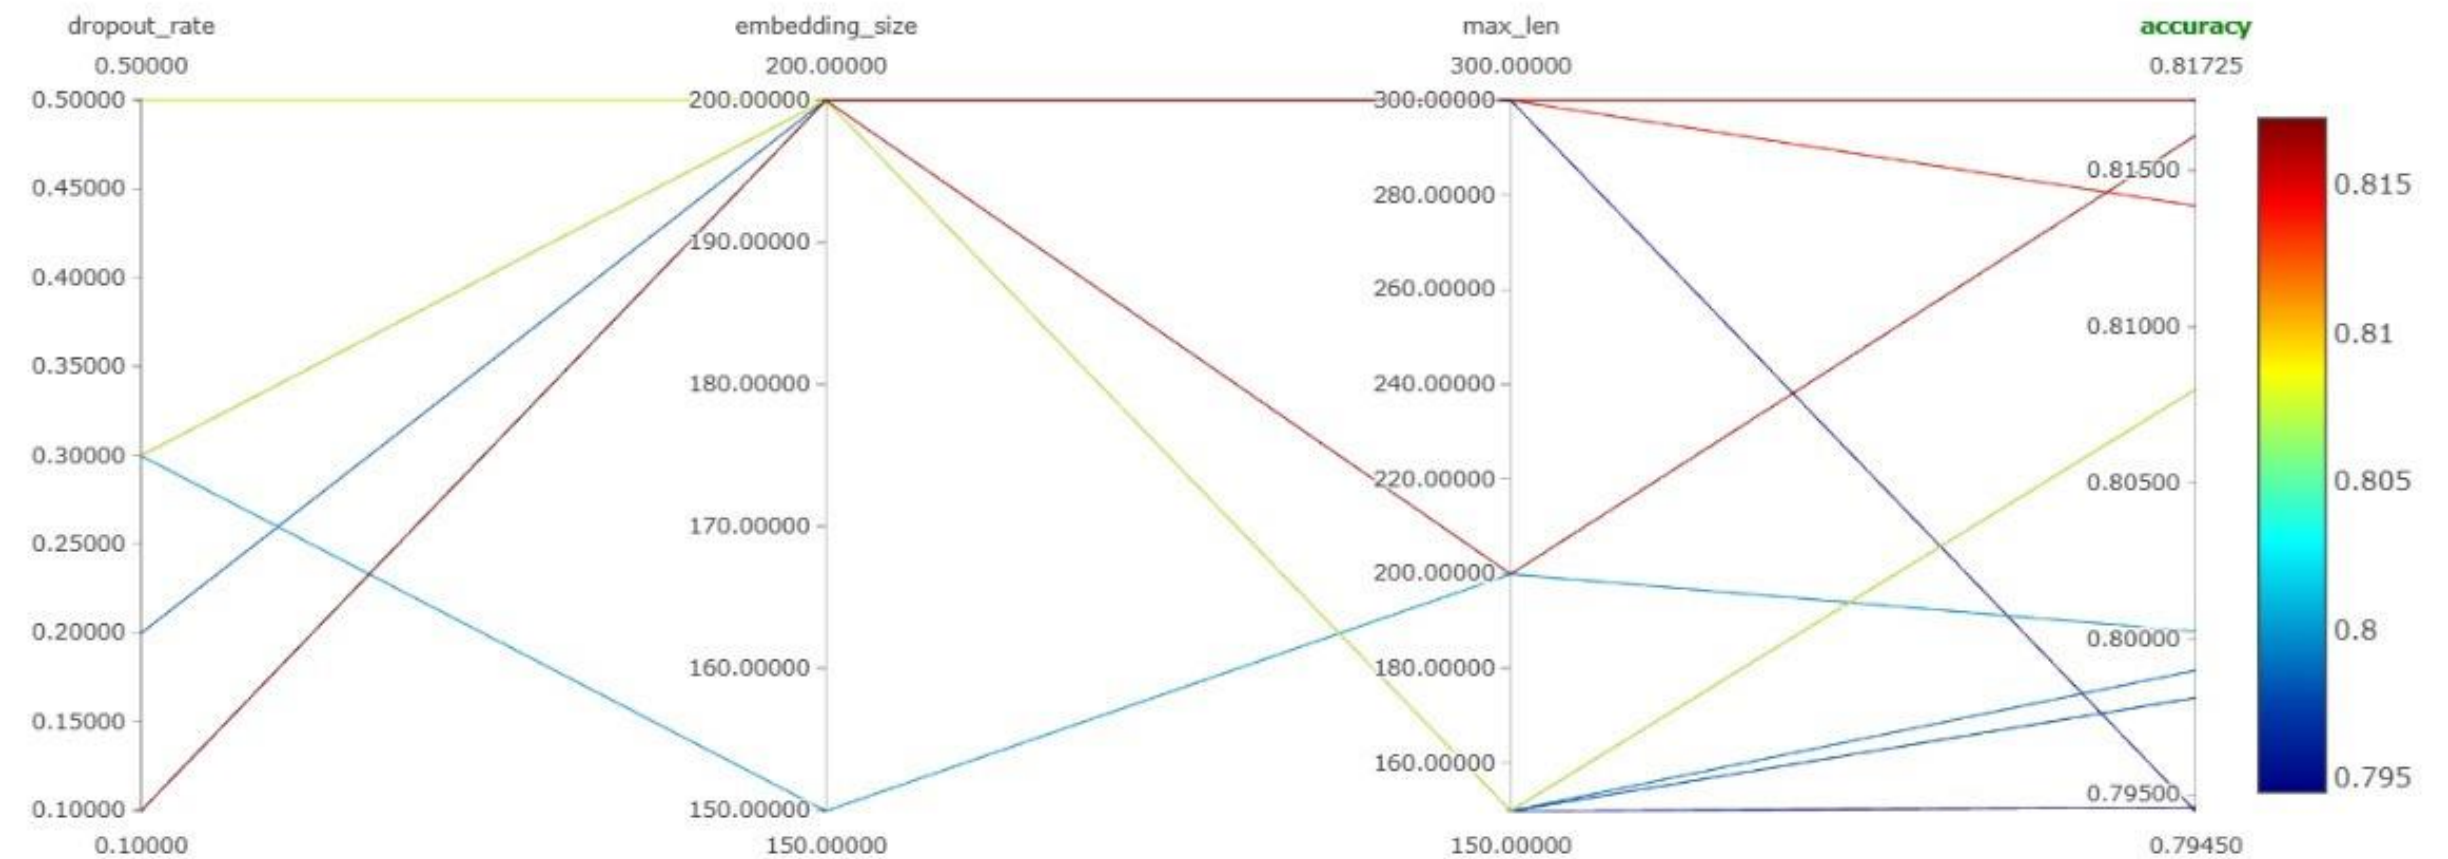
Parameters:

dropout_rate X embedding_size X max_len X

Metrics:

accuracy X

Clear All



LSTM WITH TEXT AND SUMMARY

Reviews_detection >

righteous-whale-603




Register model

Overview Model metrics System metrics **Artifacts**

- Sequential
 - LSTM_first_try
 - data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - LSTM_first_try_conf_matrix.png
 - LSTM_first_try_pr_curve.png
 - LSTM_first_try_roc_curve.png

LSTM_first_try_conf_matrix.png 11.84KB

Path: file:///c:/Users/DELL/Downloads/New%20folder%20%282%29/mlruns/940220449463371638/e5c59879170944f5bfeab26075d678f6/artifacts/LSTM_first_try_conf_ma...  




LSTM2 CONFUION MATRIX

- Sequential
 - LSTM_first_try
 - data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - LSTM_first_try_conf_matrix.png
 - LSTM_first_try_pr_curve.png
 - LSTM_first_try_roc_curve.png

LSTM_first_try_roc_curve.png 29.33KB

Path: file:///c:/Users/DELL/Downloads/New%20folder%20%282%29/mlruns/940220449463371638/2d3b0637e1ec4afdb38e163670844b53/artifacts/LSTM_first_try_roc_cur... [Download](#) [Copy](#)



LSTM2 ROC CURVE

Hugging Face



we developed a sentiment analysis project using the CardiffNLP/twitter-roberta-base-sentiment model from Hugging Face to classify reviews into positive, negative, and neutral sentiments.

Our Model

[Models](#)[Datasets](#)[Spaces](#)[Posts](#)[Docs](#)[Pricing](#)

abdelrahmanelsheikh39 / **SentimentAnalysisAtDEPI2**

like 1

Text Classification

Transformers

Safetensors

roberta

Generated from Trainer

Inference Endpoints

Model card

Files and versions

Community



Train

Deploy

Use this model

Edit model card

SentimentAnalysisAtDEPI2

This model is a fine-tuned version of [cardiffnlp/twitter-roberta-base-sentiment](#) on an unknown dataset. It achieves the following results on the evaluation set:

- Loss: 0.4876
- Accuracy: 0.8480

Model description

More information needed

Downloads last month
8



Safetensors

Model size

125M params

Tensor type

F32



Inference Examples

Text Classification

This model does not have enough activity to be deployed to Inference API (serverless) yet. Increase its social visibility and check back later, or deploy to [Inference Endpoints \(dedicated\)](#) instead.

Training Result

Training Loss	Epoch	Step	Validation Loss	Accuracy
0.5493	1.0	14212	0.5040	0.8133
0.38	2.0	28424	0.4682	0.8371
0.3531	3.0	42636	0.4678	0.8433
0.3067	4.0	56848	0.4876	0.8480

The following hyperparameters were used during training:

- learning_rate: 2e-05
- train_batch_size: 32
- eval_batch_size: 32
- seed: 42
- optimizer: Adam with betas=(0.9,0.999) and epsilon=1e-08
- lr_scheduler_type: linear
- num_epochs: 4

Training Result

```
# Use a pipeline as a high-level helper
from transformers import pipeline
```

```
pipe = pipeline("text-classification", model="cardiffnlp/twitter-roberta-base-sentiment")
pipe('I have bought several of the Vitality canned dog food products and have found them all to be of good quality. The product ]
```

```
config.json: 0%|          | 0.00/747 [00:00<?, ?B/s]
```

```
pytorch_model.bin: 0%|          | 0.00/499M [00:00<?, ?B/s]
```

```
vocab.json: 0%|          | 0.00/899k [00:00<?, ?B/s]
```

```
merges.txt: 0%|          | 0.00/456k [00:00<?, ?B/s]
```

```
special_tokens_map.json: 0%|          | 0.00/150 [00:00<?, ?B/s]
```

Hardware accelerator e.g. GPU is available in the environment, but no `device` argument is passed to the `Pipeline` object. Model will be on CPU.

```
text = "I have bought several of the Vitality canned dog food products and have found them all to be of good quality. The product
```

```
from transformers import pipeline
```

```
classifier = pipeline("text-classification", model="abdelrahmanelsheikh39/SentimentAnalysisAtDEPI2")
classifier(text)
```

```
config.json: 0%|          | 0.00/1.03k [00:00<?, ?B/s]
```

```
model.safetensors: 0%|          | 0.00/499M [00:00<?, ?B/s]
```

```
tokenizer_config.json: 0%|          | 0.00/1.24k [00:00<?, ?B/s]
```

```
vocab.json: 0%|          | 0.00/798k [00:00<?, ?B/s]
```

```
merges.txt: 0%|          | 0.00/456k [00:00<?, ?B/s]
```

```
tokenizer.json: 0%|          | 0.00/3.56M [00:00<?, ?B/s]
```

```
special_tokens_map.json: 0%|          | 0.00/958 [00:00<?, ?B/s]
```

Hardware accelerator e.g. GPU is available in the environment, but no `device` argument is passed to the `Pipeline` object. Model will be on CPU.

```
[{'label': 'LABEL_4', 'score': 0.8687160611152649}]
```


Training Result

Training results

Training Loss	Epoch	Step	Validation Loss	Accuracy
0.5583	1.0	14212	0.5413	0.7958
0.4939	2.0	28424	0.5007	0.8201
0.4564	3.0	42636	0.4923	0.8314

Training hyperparameters

The following hyperparameters were used during training:

- learning_rate: 2e-05
- train_batch_size: 32
- eval_batch_size: 32
- seed: 42
- optimizer: Adam with betas=(0.9,0.999) and epsilon=1e-08
- lr_scheduler_type: linear
- num_epochs: 3

Training Result

```
: # Use a pipeline as a high-level helper
from transformers import pipeline

pipe = pipeline("text-classification", model="nlptown/bert-base-multilingual-uncased-sentiment")
pipe("I have bought several of the Vitality canned dog food products and have found them all to be of good quality. The product is")

Hardware accelerator e.g. GPU is available in the environment, but no `device` argument is passed to the `Pipeline` object. Model will be on CPU.

: [{'label': '5 stars', 'score': 0.7185351848602295}]
```

```
text ="I have bought several of the Vitality canned dog food products and have found them all to be of good quality. The product is"
```

```
from transformers import pipeline
classifier = pipeline("text-classification", model="abdelrahmanelsheikh39/SentimentAnalysisAtDEPI")
classifier(text)
```

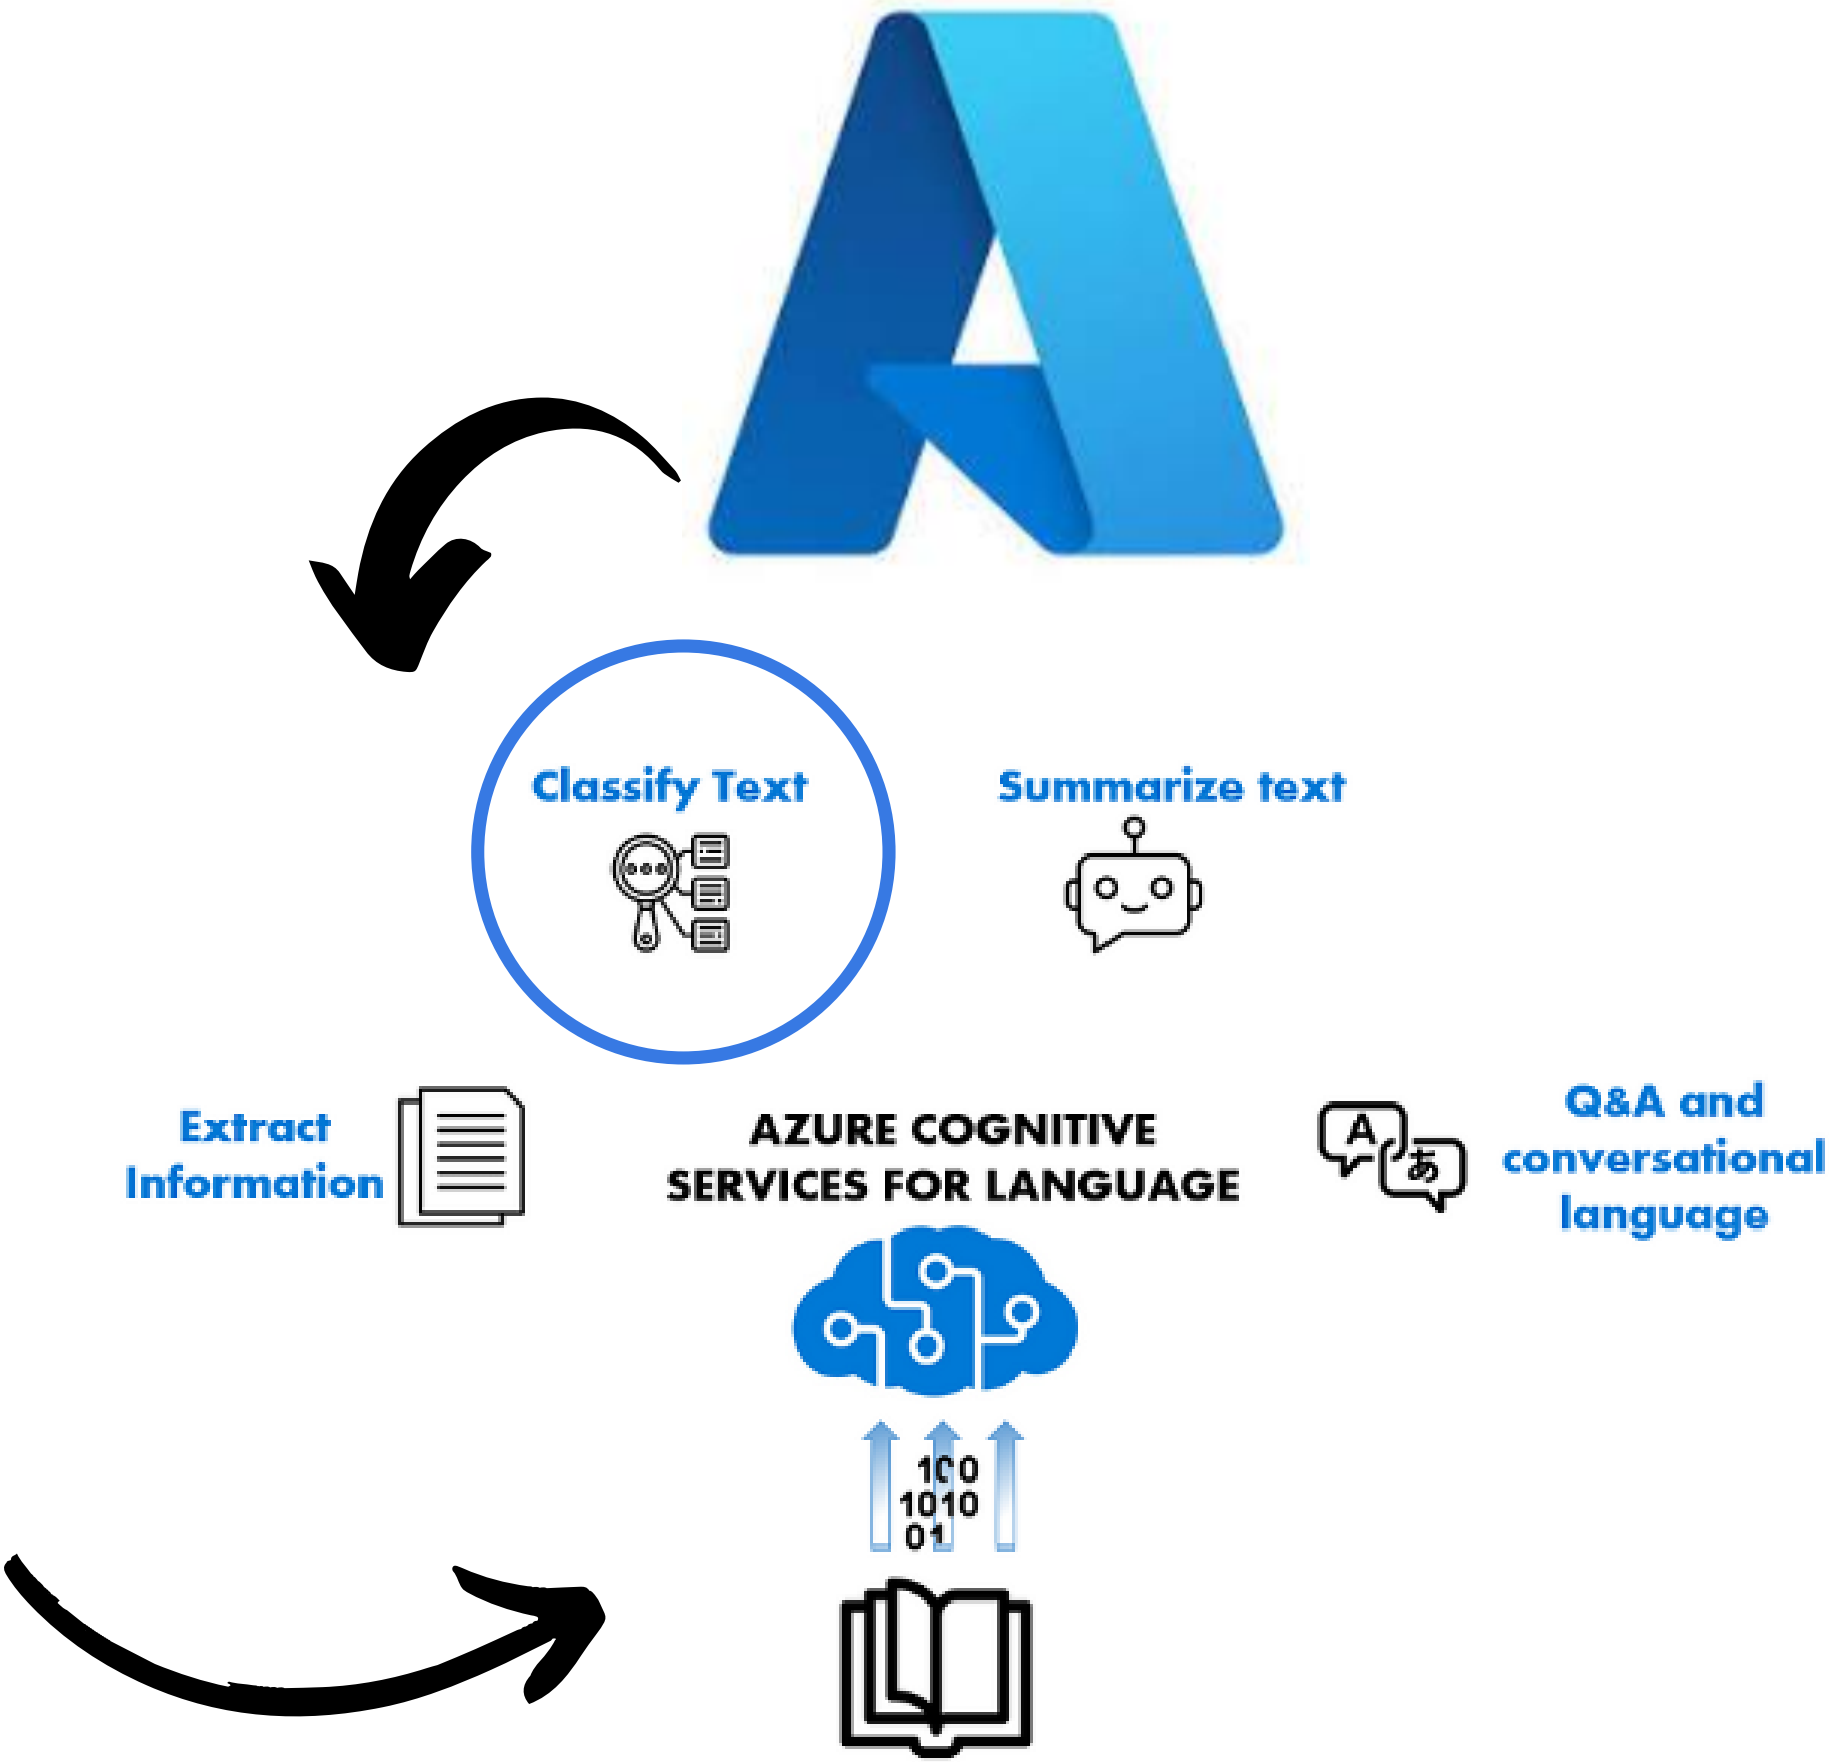
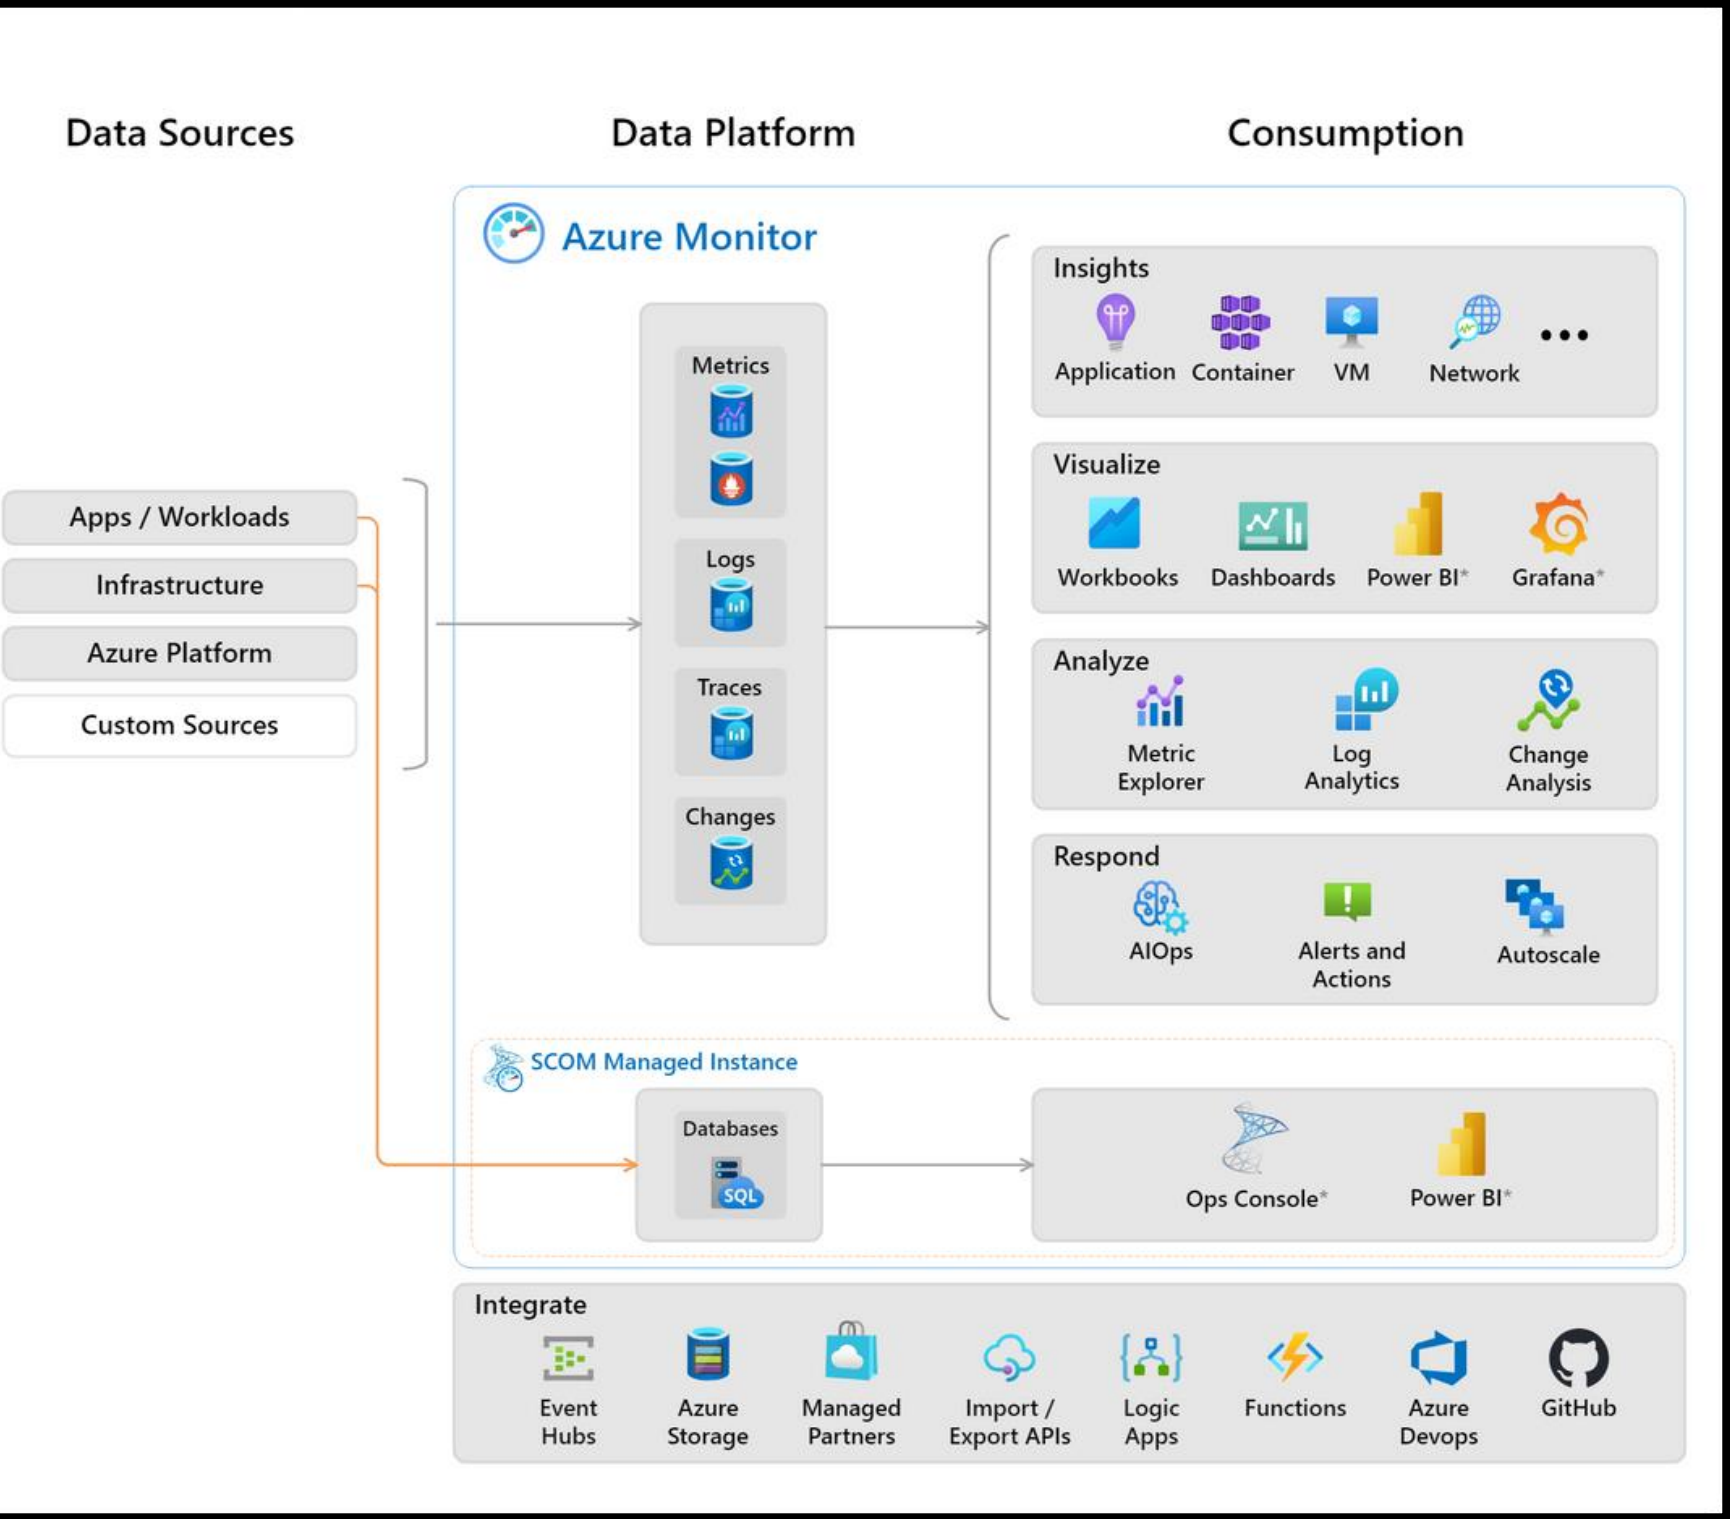
```
config.json: 0%|          | 0.00/1.23k [00:00<?, ?B/s]
model.safetensors: 0%|          | 0.00/669M [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/1.26k [00:00<?, ?B/s]
vocab.txt: 0%|          | 0.00/872k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/2.56M [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/125 [00:00<?, ?B/s]
```

Hardware accelerator e.g. GPU is available in the environment, but no `device` argument is passed to the `Pipeline` object. Model will be on CPU.

```
[{'label': '5 stars', 'score': 0.9597824215888977}]
```

Activate Windows
Go to Settings to activate Windows.


Microsoft Azure




I bought several of the Vitality canned dog food products and have found them all to be of good quality. The product looks more like a stew than a processed meat and it smells better. My Labrador is finicky and she appreciates this product better than most.

example 1


positive review



I have bought several of the Vitality canned dog food products and have found them all to be of good quality.




The product looks more like a stew than a processed meat and it smells better.



If you're impulsive like me, then \$6 is ok. Don't get me wrong, the quality of these babies is very good and I have no complaints. But in retrospect, the price is a little ridiculous (esp. when you add on the shipping).


Analyzed sentiment

Document sentiment



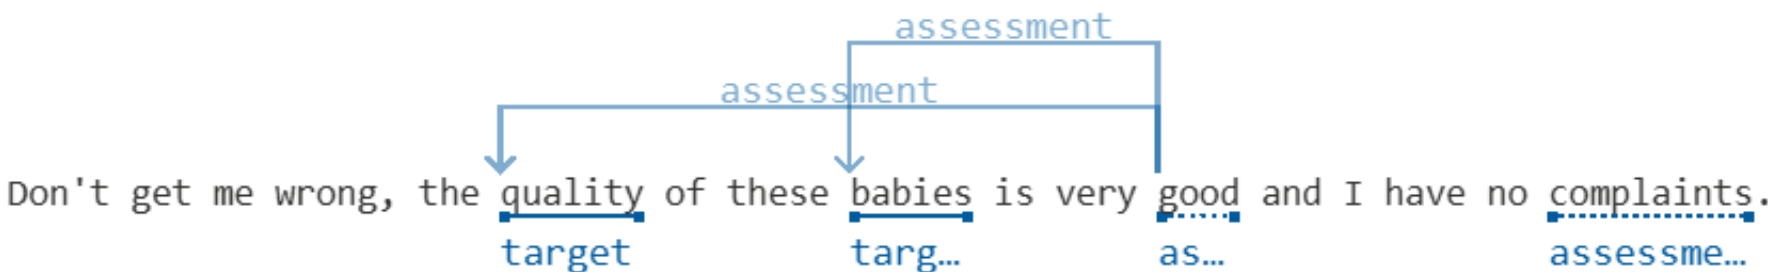
If you're impulsive like me, then \$6 is ok.

Sentence sentiment




example
2

mixed review
(mostly positive)



Sentence sentiment



Opinion

Target: quality
Assessments:
good (positive, 100.00%)

Opinion

Target: babies
Assessments:
good (positive, 100.00%)

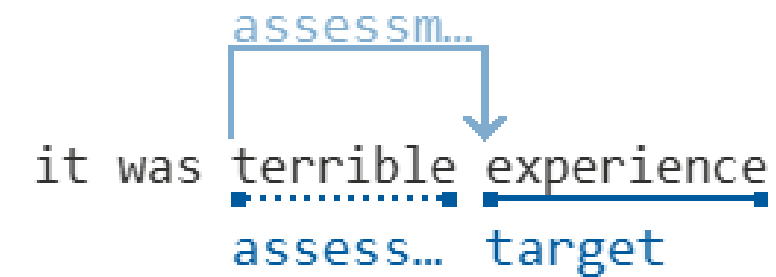
it was terrible experience

Analyzed sentiment

Document sentiment

Negative

Confidence: 0.00%



Sentence sentiment

Negative

Confidence: 0.00%



Opinion

Target: experience

Assessments:

terrible (negative,
99.00%)

example
3



Resource Page



Hugging Face



kaggle



Thank You

